

# Lab 7

Jefferson Ratliff

Math 241, Week 9

```
# Put all necessary libraries here
library(tidyverse)
library(tidytext)

# Ensure the textdata package is installed
if (!requireNamespace("textdata", quietly = TRUE)) {
  install.packages("textdata")
}
# Load the textdata package
library(textdata)

# Before knitting your document one last time, you will have to download the AFINN lexicon explicitly
lexicon_afinn()
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon      -2
## 2 abandoned    -2
## 3 abandons     -2
## 4 abducted     -2
## 5 abduction    -2
## 6 abductions    -2
## 7 abhor        -3
## 8 abhorred     -3
## 9 abhorrent    -3
## 10 abhors      -3
## # i 2,467 more rows
```

```
lexicon_nrc()
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
```

```
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # i 13,862 more rows
```

**Due: Friday, March 29th at 5:30pm**

## Goals of this lab

1. Practice matching patterns with regular expressions.
2. Practice manipulating strings with `stringr`.
3. Practice tokenizing text with `tidytext`.
4. Practice looking at word frequencies.
5. Practice conducting sentiment analysis.

### Problem 1: What's in a Name? (You'd Be Surprised!)

1. Load the `babynames` dataset, which contains yearly information on the frequency of baby names by sex and is provided by the US Social Security Administration. It includes all names with at least 5 uses per year per sex. In this problem, we are going to practice pattern matching!

```
library(babynames)
data("babynames")
#?babynames
```

- a. For 2000, find the ten most popular female baby names that start with the letter Z.

```
names_2000_Z <- filter(babynames, year == 2000 & sex == "F" & substr(name, 1, 1) == "Z")
top_names_2000_Z <- top_n(names_2000_Z, 10, n)
print(top_names_2000_Z)
```

- b. For 2000, find the ten most popular female baby names that contain the letter z.

```
names_2000_with_z <- filter(babynames, year == 2000 & sex == "F" & grepl("z", name))
top_names_2000_with_z <- head(arrange(names_2000_with_z, desc(n)), 10)
print(top_names_2000_with_z)
```

```
## # A tibble: 10 x 5
##   year sex  name      n    prop
##   <dbl> <chr> <chr>   <int>  <dbl>
## 1  2000 F    Elizabeth 15094 0.00757
## 2  2000 F    Mackenzie 6348 0.00318
## 3  2000 F    Mckenzie 2526 0.00127
## 4  2000 F    Makenzie 1613 0.000809
## 5  2000 F    Jazmin 1391 0.000697
## 6  2000 F    Jazmine 1353 0.000678
```

```
## 7 2000 F Lizbeth 817 0.000410
## 8 2000 F Eliza 759 0.000380
## 9 2000 F Litzy 722 0.000362
## 10 2000 F Esperanza 499 0.000250
```

c. For 2000, find the ten most popular female baby names that end in the letter z.

```
names_2000_end_with_z <- filter(babynames, year == 2000 & sex == "F" & substr(name, nchar(name), nchar(name)) == "z")
top_names_2000_end_with_z <- head(arrange(names_2000_end_with_z, desc(n)), 10)
print(top_names_2000_end_with_z)
```

```
## # A tibble: 10 x 5
##   year sex name      n      prop
##   <dbl> <chr> <chr>   <int>   <dbl>
## 1 2000 F Luz      489 0.000245
## 2 2000 F Beatriz 357 0.000179
## 3 2000 F Mercedes 141 0.0000707
## 4 2000 F Maricruz 96 0.0000481
## 5 2000 F Liz      72 0.0000361
## 6 2000 F Inez     69 0.0000346
## 7 2000 F Odaliz  24 0.0000120
## 8 2000 F Marycruz 23 0.0000115
## 9 2000 F Cruz    19 0.00000952
## 10 2000 F Deniz   16 0.00000802
```

d. Between your three tables in 1.a - 1.c, do any of the names show up on more than one list? If so, which ones? (Yes, I know you could do this visually but use some joins!)

```
names_2000_start_with_z <- babynames %>%
  filter(year == 2000 & sex == "F" & substr(name, 1, 1) == "Z") %>%
  select(name)

names_2000_with_z <- babynames %>%
  filter(year == 2000 & sex == "F" & grepl("z", name)) %>%
  select(name)

names_2000_end_with_z <- babynames %>%
  filter(year == 2000 & sex == "F" & substr(name, nchar(name), nchar(name)) == "z") %>%
  select(name)

common_names_1_2 <- inner_join(names_2000_start_with_z, names_2000_with_z, by = "name")
common_names_1_3 <- inner_join(names_2000_start_with_z, names_2000_end_with_z, by = "name")
common_names_2_3 <- inner_join(names_2000_with_z, names_2000_end_with_z, by = "name")

all_common_names <- bind_rows(common_names_1_2, common_names_1_3, common_names_2_3) %>%
  distinct()

# Print the common names
print(all_common_names)
```

```
## # A tibble: 41 x 1
##   name
##   <chr>
## 1 Zuzanna
## 2 Zeltzin
## 3 Zsazsa
## 4 Luz
## 5 Beatriz
## 6 Mercedes
## 7 Maricruz
## 8 Liz
## 9 Inez
## 10 Odaliz
## # i 31 more rows
```

e. Verify that none of the baby names contain a numeric (0-9) in them.

```
female_names_2000 <- babynames %>%
  filter(year == 2000 & sex == "F")

names_with_numeric <- female_names_2000 %>%
  filter(grepl("[0-9]", name))

print(names_with_numeric)
```

```
## # A tibble: 0 x 5
## # i 5 variables: year <dbl>, sex <chr>, name <chr>, n <int>, prop <dbl>
```

f. While none of the names contain 0-9, that doesn't mean they don't contain "one", "two", ..., or "nine". Create a table that provides the number of times a baby's name contained the word "zero", the word "one", ... the word "nine".

```
female_names_2000 <- babynames %>%
  filter(year == 2000 & sex == "F")

number_words <- c("zero", "one", "two", "three", "four", "five", "six", "seven", "eight", "nine")

number_counts <- data.frame(Number = number_words, Count = numeric(length(number_words)))

for (i in 1:length(number_words)) {
  word <- number_words[i]
  counts <- female_names_2000 %>%
    filter(grepl(paste0("\\b", word, "\\b"), name)) %>%
    summarise(count = n())
  number_counts[i, "Count"] <- counts$count
}

print(number_counts)
```

```
##   Number Count
## 1    zero    0
```

```
## 2    one    0
## 3    two    0
## 4   three    0
## 5    four    0
## 6    five    0
## 7    six     0
## 8   seven    0
## 9   eight    0
## 10   nine    0
```

Notes:

- I recommend first converting all the names to lower case.
- If none of the baby's names contain the written number, there you can leave the number out of the table.
- Use `str_extract()`, not `str_extract_all()`. (We will ignore names where more than one of the words exists.)

*Hint:* You will have two steps that require pattern matching: 1. Subset your table to only include the rows with the desired words. 2. Add a column that contains the desired word.

g. Which written number or numbers don't show up in any of the baby names?

```
female_names_2000 <- babynames %>%
  filter(year == 2000 & sex == "F") %>%
  mutate(name = tolower(name))

number_words <- c("zero", "one", "two", "three", "four", "five", "six", "seven", "eight", "nine")

number_counts <- data.frame(Number = character(), Count = integer(), stringsAsFactors = FALSE)

for (word in number_words) {

  names_with_word <- female_names_2000 %>%
    filter(str_detect(name, paste0("\\b", word, "\\b")))

  names_with_word <- names_with_word %>%
    mutate(word = str_extract(name, paste0("\\b", word, "\\b"))) %>%
    distinct() # Keep only unique names

  word_count <- nrow(names_with_word)

  if (word_count > 0) {
    number_counts <- bind_rows(number_counts, data.frame(Number = word, Count = word_count))
  }
}

print(number_counts)
```

```
##   Number Count
## 1   seven     1
```

```
missing_numbers <- setdiff(number_words, number_counts$Number)
print(paste("Written number(s) not found in any baby names:", paste(missing_numbers, collapse = ", ")))
```

```
## [1] "Written number(s) not found in any baby names: zero, one, two, three, four, five, six, eight, n
```

h. Create a table that contains the names and their frequencies for the two least common written numbers.

```
female_names_2000 <- babynames %>%
  filter(year == 2000 & sex == "F") %>%
  mutate(name = tolower(name))

number_words <- c("zero", "one", "two", "three", "four", "five", "six", "seven", "eight", "nine")

number_counts <- data.frame(Number = character(), Count = integer(), stringsAsFactors = FALSE)

for (word in number_words) {
  names_with_word <- female_names_2000 %>%
    filter(str_detect(name, paste0("\\b", word, "\\b")))

  names_with_word <- names_with_word %>%
    mutate(word = str_extract(name, paste0("\\b", word, "\\b"))) %>%
    distinct()

  word_count <- nrow(names_with_word)

  if (word_count > 0) {
    number_counts <- bind_rows(number_counts, data.frame(Number = word, Count = word_count))
  }
}

least_common_numbers <- number_counts %>%
  arrange(Count) %>%
  head(2)

names_least_common_numbers <- female_names_2000 %>%
  filter(str_detect(name, paste0("\\b", least_common_numbers$Number[1], "\\b|\\b", least_common_numbers$Number[2], "\\b")))
  group_by(name) %>%
  summarise(total_count = sum(n))

print(names_least_common_numbers)
```

```
## # A tibble: 1 x 2
##   name total_count
##   <chr>      <int>
## 1 seven         24
```

i. List out the names that contain no vowels (consider “y” to be a vowel).

```
female_names_2000 <- babynames %>%
  filter(year == 2000 & sex == "F")

contains_no_vowels <- function(name) {
  !grepl("[aeiouyAEIOUY]", name)
}

names_no_vowels <- female_names_2000 %>%
  filter(sapply(name, contains_no_vowels))

print(names_no_vowels$name)
```

```
## [1] "Bg" "Kc" "Mc"
```

## Problem 2: Tidying the “Call of the Wild”

Did you read “Call of the Wild” by Jack London? If not, [read the first paragraph of its wiki page](#) for a quick summary and then let’s do some text analysis on this classic! The following code will pull the book into R using the `gutenbergr` package.

```
library(gutenbergr)
wild <- gutenberg_download(215)
```

- a. Create a tidy text dataset where you tokenize by words.

```
wild_tokens <- wild %>%
  unnest_tokens(word, text)

head(wild_tokens)
```

```
## # A tibble: 6 x 2
##   gutenberg_id word
##         <int> <chr>
## 1         215 cover
## 2         215 the
## 3         215 call
## 4         215 of
## 5         215 the
## 6         215 wild
```

- b. Find the frequency of the 20 most common words. First, remove stop words.

```
wild_tokens <- wild %>%
  unnest_tokens(word, text)

wild_tokens_no_stop <- wild_tokens %>%
  anti_join(stop_words)

word_counts <- wild_tokens_no_stop %>%
  count(word, sort = TRUE)
```

```
top_20_words <- head(word_counts, 20)

print(top_20_words)
```

```
## # A tibble: 20 x 2
##   word      n
##   <chr>   <int>
## 1 buck    313
## 2 dogs    118
## 3 thornton 81
## 4 dog     79
## 5 time    77
## 6 day     70
## 7 life    64
## 8 sled    63
## 9 half    60
## 10 spitz   60
## 11 head    54
## 12 camp    53
## 13 françois 51
## 14 feet    49
## 15 buck's  47
## 16 trail   42
## 17 days    41
## 18 eyes    40
## 19 john    40
## 20 night   40
```

c. Create a bar graph and a word cloud of the frequencies of the 20 most common words.

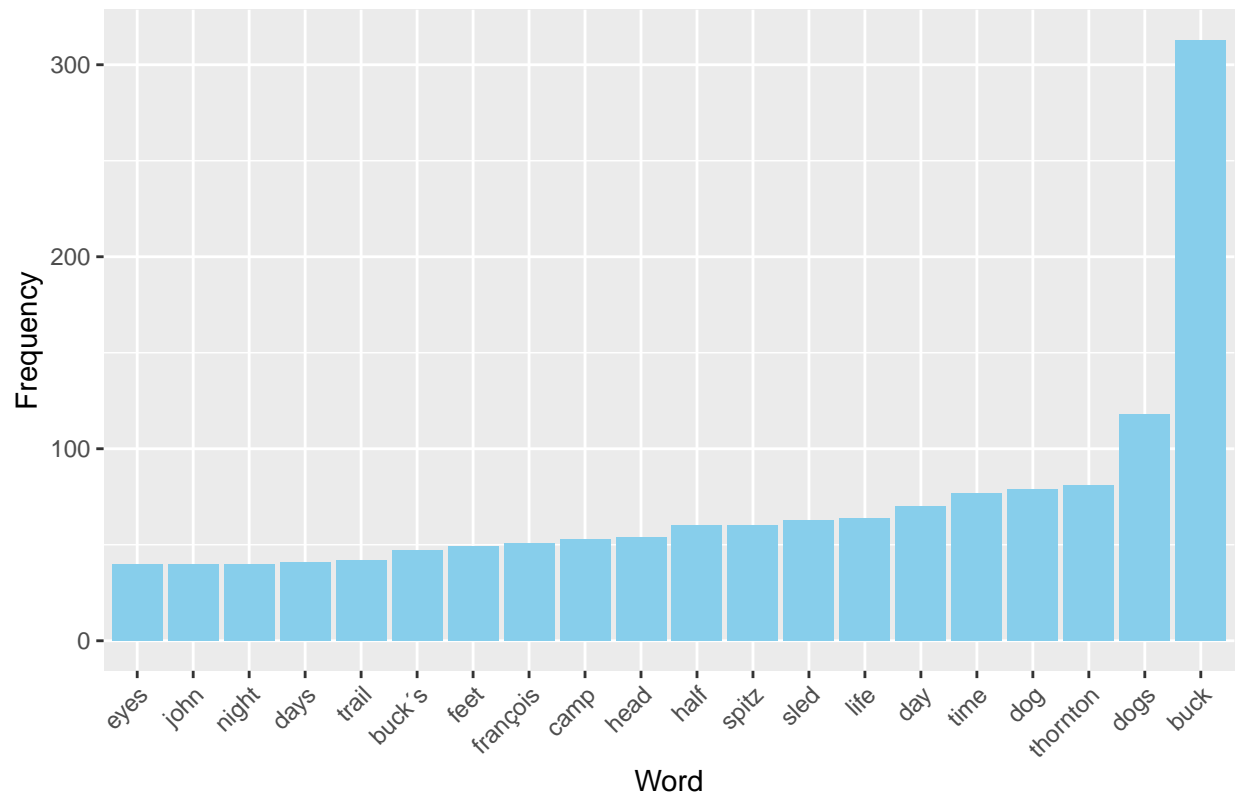
```
library(wordcloud)

bar_plot <- ggplot(top_20_words, aes(x = reorder(word, n), y = n)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(x = "Word", y = "Frequency", title = "Top 20 Most Common Words in 'Call of the Wild'") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

print(bar_plot)
```



Top 20 Most Common Words in 'Call of the Wild'



```
wordcloud(words = top_20_words$word, freq = top_20_words$n, scale=c(3,0.5), min.freq = 1,
          random.order=FALSE, rot.per=0.35, colors=brewer.pal(8, "Dark2"))
```



- d. Explore the sentiment of the text using three of the sentiment lexicons in `tidytext`. What does your analysis say about the sentiment of the text?

Notes:

- Make sure to NOT remove stop words this time.
- `afinn` is a numeric score and should be handled differently than the categorical scores.

```
sentiment_bing <- wild_tokens %>%
  inner_join(get_sentiments("bing"))

sentiment_afinn <- wild_tokens %>%
  inner_join(get_sentiments("afinn"))

sentiment_nrc <- wild_tokens %>%
  inner_join(get_sentiments("nrc"))

bing_score <- sum(sentiment_bing$sentiment == "positive") - sum(sentiment_bing$sentiment == "negative")
afinn_score <- sum(sentiment_afinn$value)
nrc_score <- sum(sentiment_nrc$sentiment == "positive") - sum(sentiment_nrc$sentiment == "negative")

print(paste("Bing sentiment score:", bing_score))
```

```
## [1] "Bing sentiment score: -675"
```

```
print(paste("AFINN sentiment score:", afinn_score))
```

```
## [1] "AFINN sentiment score: -591"
```

```
print(paste("NRC sentiment score:", nrc_score))
```

```
## [1] "NRC sentiment score: -266"
```

- e. If you didn't do so in 2.d, compute the average sentiment score of the text using `afinn`. Which positive words had the biggest impact? Which negative words had the biggest impact?

```
sentiment_afinn %>%  
  summarize(mean = mean(value))
```

```
## # A tibble: 1 x 1  
##   mean  
##   <dbl>  
## 1 -0.322
```

```
sentiment_afinn %>%  
  group_by(word) %>%  
  summarize(mean = mean(value),  
            n = n()) %>%  
  arrange(desc(n))
```

```
## # A tibble: 528 x 3  
##   word      mean      n  
##   <chr>    <dbl> <int>  
## 1 no        -1     95  
## 2 like        2     58  
## 3 great        3     49  
## 4 fire       -2     33  
## 5 good         3     28  
## 6 dead       -3     22  
## 7 cried      -2     21  
## 8 love         3     21  
## 9 strength    2     21  
## 10 broke     -1     18  
## # i 518 more rows
```

- f. You should have found that “no” was an important negative word in the sentiment score. To know if that really makes sense, let's turn to the raw lines of text for context. Pull out all of the lines that have the word “no” in them. Make sure to not pull out extraneous lines (e.g., a line with the word “now”).

```
no_wild <- wild %>%  
  filter(str_detect(text, "\\bno\\b"))  
no_wild
```

```
## # A tibble: 83 x 2
##   gutenbergs_id text
##   <int> <chr>
## 1      215 solitary man, no one saw them arrive at the little flag station~
## 2      215 that it was the club, but his madness knew no caution. A dozen ~
## 3      215 "He's no slouch at dog-breakin', that's wot I say," one of the ~
## 4      215 all, that he stood no chance against a man with a club. He had ~
## 5      215 in the red sweater. "And seem' it's government money, you ain't~
## 6      215 animal. The Canadian Government would be no loser, nor would its
## 7      215 while he developed no affection for them, he none the less grew
## 8      215 The other dog made no advances, nor received any; also, he did ~
## 9      215 knew no law but the law of club and fang.
## 10     215 full-grown wolf, though not half so large as she. There was no ~
## # i 73 more rows
```

g. Draw some conclusions about how “no” is used in the text.

No is used not a lot within character dialogue but more as a negative modifier to communicate the lack of a value or thing.

h. We can also look at how the sentiment of the text changes as the text progresses. Below, I have added two columns to the original dataset. Now I want you to do the following wrangling:

- Tidy the data (but don't drop stop words).
- Add the word sentiments using `bing`.
- Count the frequency of sentiments by index.
- Reshape the data to be wide with the count of the negative sentiments in one column and the positive in another, along with a column for index.
- Compute a sentiment column by subtracting the negative score from the positive.

```
wild_time <- wild %>%
  mutate(line = row_number(), index = floor(line/45) + 1)
```

```
#Hint: fill = 0 will insert zero instead of NA
pivot_XXX(..., values_fill = 0)
```

```
wild_tidy <- wild_time %>%
  unnest_tokens(word, text)

wild_sentiments <- wild_tidy %>%
  inner_join(get_sentiments("bing"))

sentiment_counts <- wild_sentiments %>%
  count(index, sentiment)

sentiment_wide <- sentiment_counts %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0)

sentiment_wide <- sentiment_wide %>%
  mutate(sentiment = positive - negative)


print(sentiment_wide)
```

```
## # A tibble: 69 x 4
##   index negative positive sentiment
##   <dbl>   <int>   <int>   <int>
## 1     1       9       7       -2
## 2     2       9      15        6
## 3     3      26      17       -9
## 4     4      17       6      -11
## 5     5      29      10      -19
## 6     6      25       8      -17
## 7     7      22      14       -8
## 8     8      12      11        -1
## 9     9      20      19        -1
## 10    10      24      11      -13
## # i 59 more rows
```

i. Create a plot of the sentiment scores as the text progresses.

```
sentiment_plot <- ggplot(sentiment_wide, aes(x = index, y = sentiment)) +
  geom_line(color = "blue") +
  geom_point(color = "blue", size = 2) +
  labs(x = "Index", y = "Sentiment Score", title = "Sentiment Score Progression") +
  theme_minimal()

# Print the plot
print(sentiment_plot)
```



- j. The choice of 45 lines per chunk was pretty arbitrary. Try modifying the index value a few times and recreating the plot in i. Based on your plots, what can you conclude about the sentiment of the novel as it progresses?

```
wild_time_modified <- wild %>%
  mutate(line = row_number(), index = floor(line/30) + 1)

wild_tidy_modified <- wild_time_modified %>%
  unnest_tokens(word, text)

wild_sentiments_modified <- wild_tidy_modified %>%
  inner_join(get_sentiments("bing"))

sentiment_counts_modified <- wild_sentiments_modified %>%
  count(index, sentiment)


sentiment_wide_modified <- sentiment_counts_modified %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0)

sentiment_wide_modified <- sentiment_wide_modified %>%
  mutate(sentiment = positive - negative)

sentiment_plot_modified <- ggplot(sentiment_wide_modified, aes(x = index, y = sentiment)) +
  geom_line(color = "blue") +
  geom_point(color = "blue", size = 2) +
  labs(x = "Index", y = "Sentiment Score", title = "Sentiment Score Progression (Modified Index)") +
  theme_minimal()

print(sentiment_plot_modified)
```

Sentiment Score Progression (Modified Index)



k. Let's look at the bigrams (2 consecutive words). Tokenize the text by bigrams.

```
wild_tokens <- wild %>%
  unnest_tokens(word, text)

wild_bigrams <- wild_tokens %>%
  mutate(next_word = lead(word)) %>%
  filter(!is.na(next_word)) %>%
  unite(bigram, word, next_word, sep = " ")

head(wild_bigrams)
```

```
## # A tibble: 6 x 2
##   gutenber_id bigram
##   <int> <chr>
## 1      215 cover the
## 2      215 the call
## 3      215 call of
## 4      215 of the
## 5      215 the wild
## 6      215 wild by
```

l. Produce a sorted table that counts the frequency of each bigram and notice that stop words are still an issue.

```
bigram_freq <- wild_bigrams %>%  
  count(bigram)
```

```
bigram_freq <- bigram_freq %>%  
  arrange(desc(n))  
bigram_freq
```

```
## # A tibble: 20,538 x 2  
##   bigram      n  
##   <chr>    <int>  
## 1 of the    246  
## 2 in the    176  
## 3 he was    131  
## 4 to the    122  
## 5 it was    111  
## 6 and the   103  
## 7 on the     86  
## 8 he had     83  
## 9 at the     71  
## 10 into the  69  
## # i 20,528 more rows
```